

```
In [ ]: 1 # !pip install -U kaleido
```

```
In [ ]: 1 # !pip install --upgrade "kaleido==0.1.*"
```

Looking in indexes: <https://pypi.org/simple>, (<https://pypi.org/simple>,) <https://us-python.pkg.dev/colab-wheels/public/simple/> (<https://us-python.pkg.dev/colab-wheels/public/simple/>)
Requirement already satisfied: kaleido==0.1.* in /usr/local/lib/python3.7/dist-packages (0.1.0)

```
In [135]: 1 # importing all the necessary libraries/modules
2
3 import io
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 import seaborn as sns
8 import random
9 import json
10 import math
11 import kaleido
12 from sklearn.datasets import load_boston
13 from sklearn.preprocessing import LabelEncoder
14 from sklearn.preprocessing import OneHotEncoder
15 from sklearn.preprocessing import StandardScaler
16 from sklearn.preprocessing import MinMaxScaler
17 from sklearn.linear_model import LinearRegression
18 from sklearn.model_selection import train_test_split
19 from sklearn.metrics import r2_score
20 from sklearn.metrics import mean_absolute_error
21 from sklearn.metrics import mean_squared_error
22 from sklearn.metrics import classification_report
23 from yellowbrick.regressor import ResidualsPlot
24 from sklearn.linear_model import Ridge
25 from sklearn.tree import DecisionTreeRegressor
26 from sklearn.tree import plot_tree
27 import matplotlib.pyplot as plt
28 from sklearn.ensemble import RandomForestRegressor
29 from sklearn.linear_model import SGDRegressor
30 from sklearn.model_selection import cross_val_score
31 from sklearn.preprocessing import scale
32 from sklearn.compose import TransformedTargetRegressor
33 from sklearn.preprocessing import QuantileTransformer
34 from sklearn.ensemble import GradientBoostingRegressor
35 import lightgbm
36 import xgboost
```

Importing the dataset:

```
In [2]: 1 # Importing the dataset onto python
2 df = pd.read_csv('Life Expectancy Data_HV22.csv')
```

```
In [3]: 1 # Moving the Target column to the end
2
3 def moveColumn(df, column_name, position):
4     column_to_move = df.pop(column_name)
5     df.insert(position, column_name, column_to_move)
6
7 moveColumn(df, 'Life expectancy ', 21)
```

```
In [4]: 1 df.head()
```

Out[4]:

	Country	Year	Status	Adult Mortality	infant deaths	Alcohol	percentage expenditure	Hepatitis B	Measles	BMI	...	Total expenditure	Diphtheria	HIV/AIDS	GDP	Population	thinness 1-19 years	thinness 5-9 years	Income composition of resources	Schooling	Life expectancy
0	Afghanistan	2015	Developing	263.0	62	0.01	71.279624	65.0	1154	19.1	...	8.16	65.0	0.1	584.259210	33736494.0	17.2	17.3	0.479	10.1	65.0
1	Afghanistan	2014	Developing	271.0	64	0.01	73.523582	62.0	492	18.6	...	8.18	62.0	0.1	612.696514	327582.0	17.5	17.5	0.476	10.0	59.9
2	Afghanistan	2013	Developing	268.0	66	0.01	73.219243	64.0	430	18.1	...	8.13	64.0	0.1	631.744976	31731688.0	17.7	17.7	0.470	9.9	59.9
3	Afghanistan	2012	Developing	272.0	69	0.01	78.184215	67.0	2787	17.6	...	8.52	67.0	0.1	669.959000	3696958.0	17.9	18.0	0.463	9.8	59.5
4	Afghanistan	2011	Developing	275.0	71	0.01	7.097109	68.0	3013	17.2	...	7.87	68.0	0.1	63.537231	2978599.0	18.2	18.2	0.454	9.5	59.2

5 rows × 22 columns

Basic Summary about df:

```
In [5]: 1 # Basic statistical summary of the df
2 df.describe()
```

Out[5]:

	Year	Adult Mortality	infant deaths	Alcohol	percentage expenditure	Hepatitis B	Measles	BMI	under-five deaths	Polio	Total expenditure	Diphtheria	HIV/AIDS	GDP	Population	thinness 1-19 years	thinness 5-9 years	co
count	2938.000000	2928.000000	2938.000000	2744.000000	2938.000000	2385.000000	2938.000000	2904.000000	2938.000000	2919.000000	2712.00000	2919.000000	2938.000000	2490.000000	2.286000e+03	2904.000000	2904.000000	27
mean	2007.518720	164.796448	30.303948	4.602861	738.251295	80.940461	2419.592240	38.321247	42.035739	82.550188	5.93819	82.324084	1.742103	7483.158469	1.275338e+07	4.839704	4.870317	
std	4.613841	124.292079	117.926501	4.052413	1987.914858	25.070016	11467.272489	20.044034	160.445548	23.428046	2.49832	23.716912	5.077785	14270.169342	6.101210e+07	4.420195	4.508882	
min	2000.000000	1.000000	0.000000	0.010000	0.000000	1.000000	0.000000	1.000000	0.000000	3.000000	0.37000	2.000000	0.100000	1.681350	3.400000e+01	0.100000	0.100000	
25%	2004.000000	74.000000	0.000000	0.877500	4.685343	77.000000	0.000000	19.300000	0.000000	78.000000	4.26000	78.000000	0.100000	463.935626	1.957932e+05	1.600000	1.500000	
50%	2008.000000	144.000000	3.000000	3.755000	64.912906	92.000000	17.000000	43.500000	4.000000	93.000000	5.75500	93.000000	0.100000	1766.947595	1.386542e+06	3.300000	3.300000	
75%	2012.000000	228.000000	22.000000	7.702500	441.534144	97.000000	360.250000	56.200000	28.000000	97.000000	7.49250	97.000000	0.800000	5910.806335	7.420359e+06	7.200000	7.200000	
max	2015.000000	723.000000	1800.000000	17.870000	19479.911610	99.000000	212183.000000	87.300000	2500.000000	99.000000	17.60000	99.000000	50.600000	119172.741800	1.293859e+09	27.700000	28.600000	

```
In [6]: 1 # data types of all the columns
        2 df.dtypes
```

Out[6]: Country object
Year int64
Status object
Adult Mortality float64
infant deaths int64
Alcohol float64
percentage expenditure float64
Hepatitis B float64
Measles int64
BMI float64
under-five deaths int64
Polio float64
Total expenditure float64
Diphtheria float64
HIV/AIDS float64
GDP float64
Population float64
thinness 1-19 years float64
thinness 5-9 years float64
Income composition of resources float64
Schooling float64
Life expectancy float64
dtype: object

```
In [7]: 1 # Prnitin the basic idea about the df
        2 print("Total number of NA values in df:", df.isna().sum().sum(), "\nTotal number of columns in df:", len(df.columns), "\nThere are", len(df['Country'].unique()), "countries in total.")
```

Total number of NA values in df: 2563
Total number of columns in df: 22
There are 193 countries in total.

Dealing with all the missing values:

```
In [8]: 1 # Calculating the number of NA values in each column
        2 df.isna().sum()
```

```
Out[8]: Country      0
        Year      0
        Status      0
        Adult Mortality  10
        infant deaths  0
        Alcohol    194
        percentage expenditure  0
        Hepatitis B  553
        Measles     0
         BMI       34
        under-five deaths  0
        Polio      19
        Total expenditure  226
        Diphtheria  19
         HIV/AIDS   0
        GDP       448
        Population  652
         thinness  1-19 years  34
         thinness 5-9 years  34
        Income composition of resources  167
        Schooling   163
        Life expectancy  10
        dtype: int64
```

Country

```
In [9]: 1 # Eliminating all the countries with only 1 instance
        2 countries = df["Country"].value_counts()
        3 oneCountry = list(countries[countries == 1].index)
        4 oneCountry
```

```
Out[9]: ['Tuvalu',
        'Cook Islands',
        'Marshall Islands',
        'Monaco',
        'Palau',
        'Niue',
        'San Marino',
        'Nauru',
        'Saint Kitts and Nevis',
        'Dominica']
```


```
In [10]: 1 for element in df["Country"]:
          2     if element in oneCountry:
          3         indices = df[df["Country"]==element].index
          4         df.drop(indices,axis=0, inplace=True)
```

```
In [11]: 1 # resetting the index of the df
          2 df.reset_index(drop=True, inplace=True)
```

Rest of the columns

```
In [12]: 1 # Plotting a correlation heatmap for all the NaN values
2 sns.heatmap(df.isnull(),yticklabels=False,cbar=False,cmap="viridis")
```

Out[12]: <AxesSubplot:>



```
In [13]: 1 # Filling in all the NaN values through the 'interpolate' method
2
3 df.reset_index(inplace=True)
4 df.groupby('Country').apply(lambda group: group.interpolate(method= 'linear'))
5 imputed_data = []
6 for year in list(df.Year.unique()):
7     year_data = df[df.Year == year].copy()
8     for col in list(year_data.columns)[4:]:
9         year_data[col] = year_data[col].fillna(year_data[col].dropna().median()).copy()
10    imputed_data.append(year_data)
11 df = pd.concat(imputed_data).copy()
```

<ipython-input-13-e45010ce39de>:4: FutureWarning: Not prepending group keys to the result index of transform-like apply. In the future, the group keys will be included in the index, regardless of whether the applied function returns a like-indexed object.
To preserve the previous behavior, use

```
>>> .groupby(..., group_keys=False)
```

To adopt the future behavior and silence this warning, use

```
>>> .groupby(..., group_keys=True)
df.groupby('Country').apply(lambda group: group.interpolate(method= 'linear'))
```

```
In [14]: 1 print("The number of NaN values in the df:", df.isna().sum().sum())
```

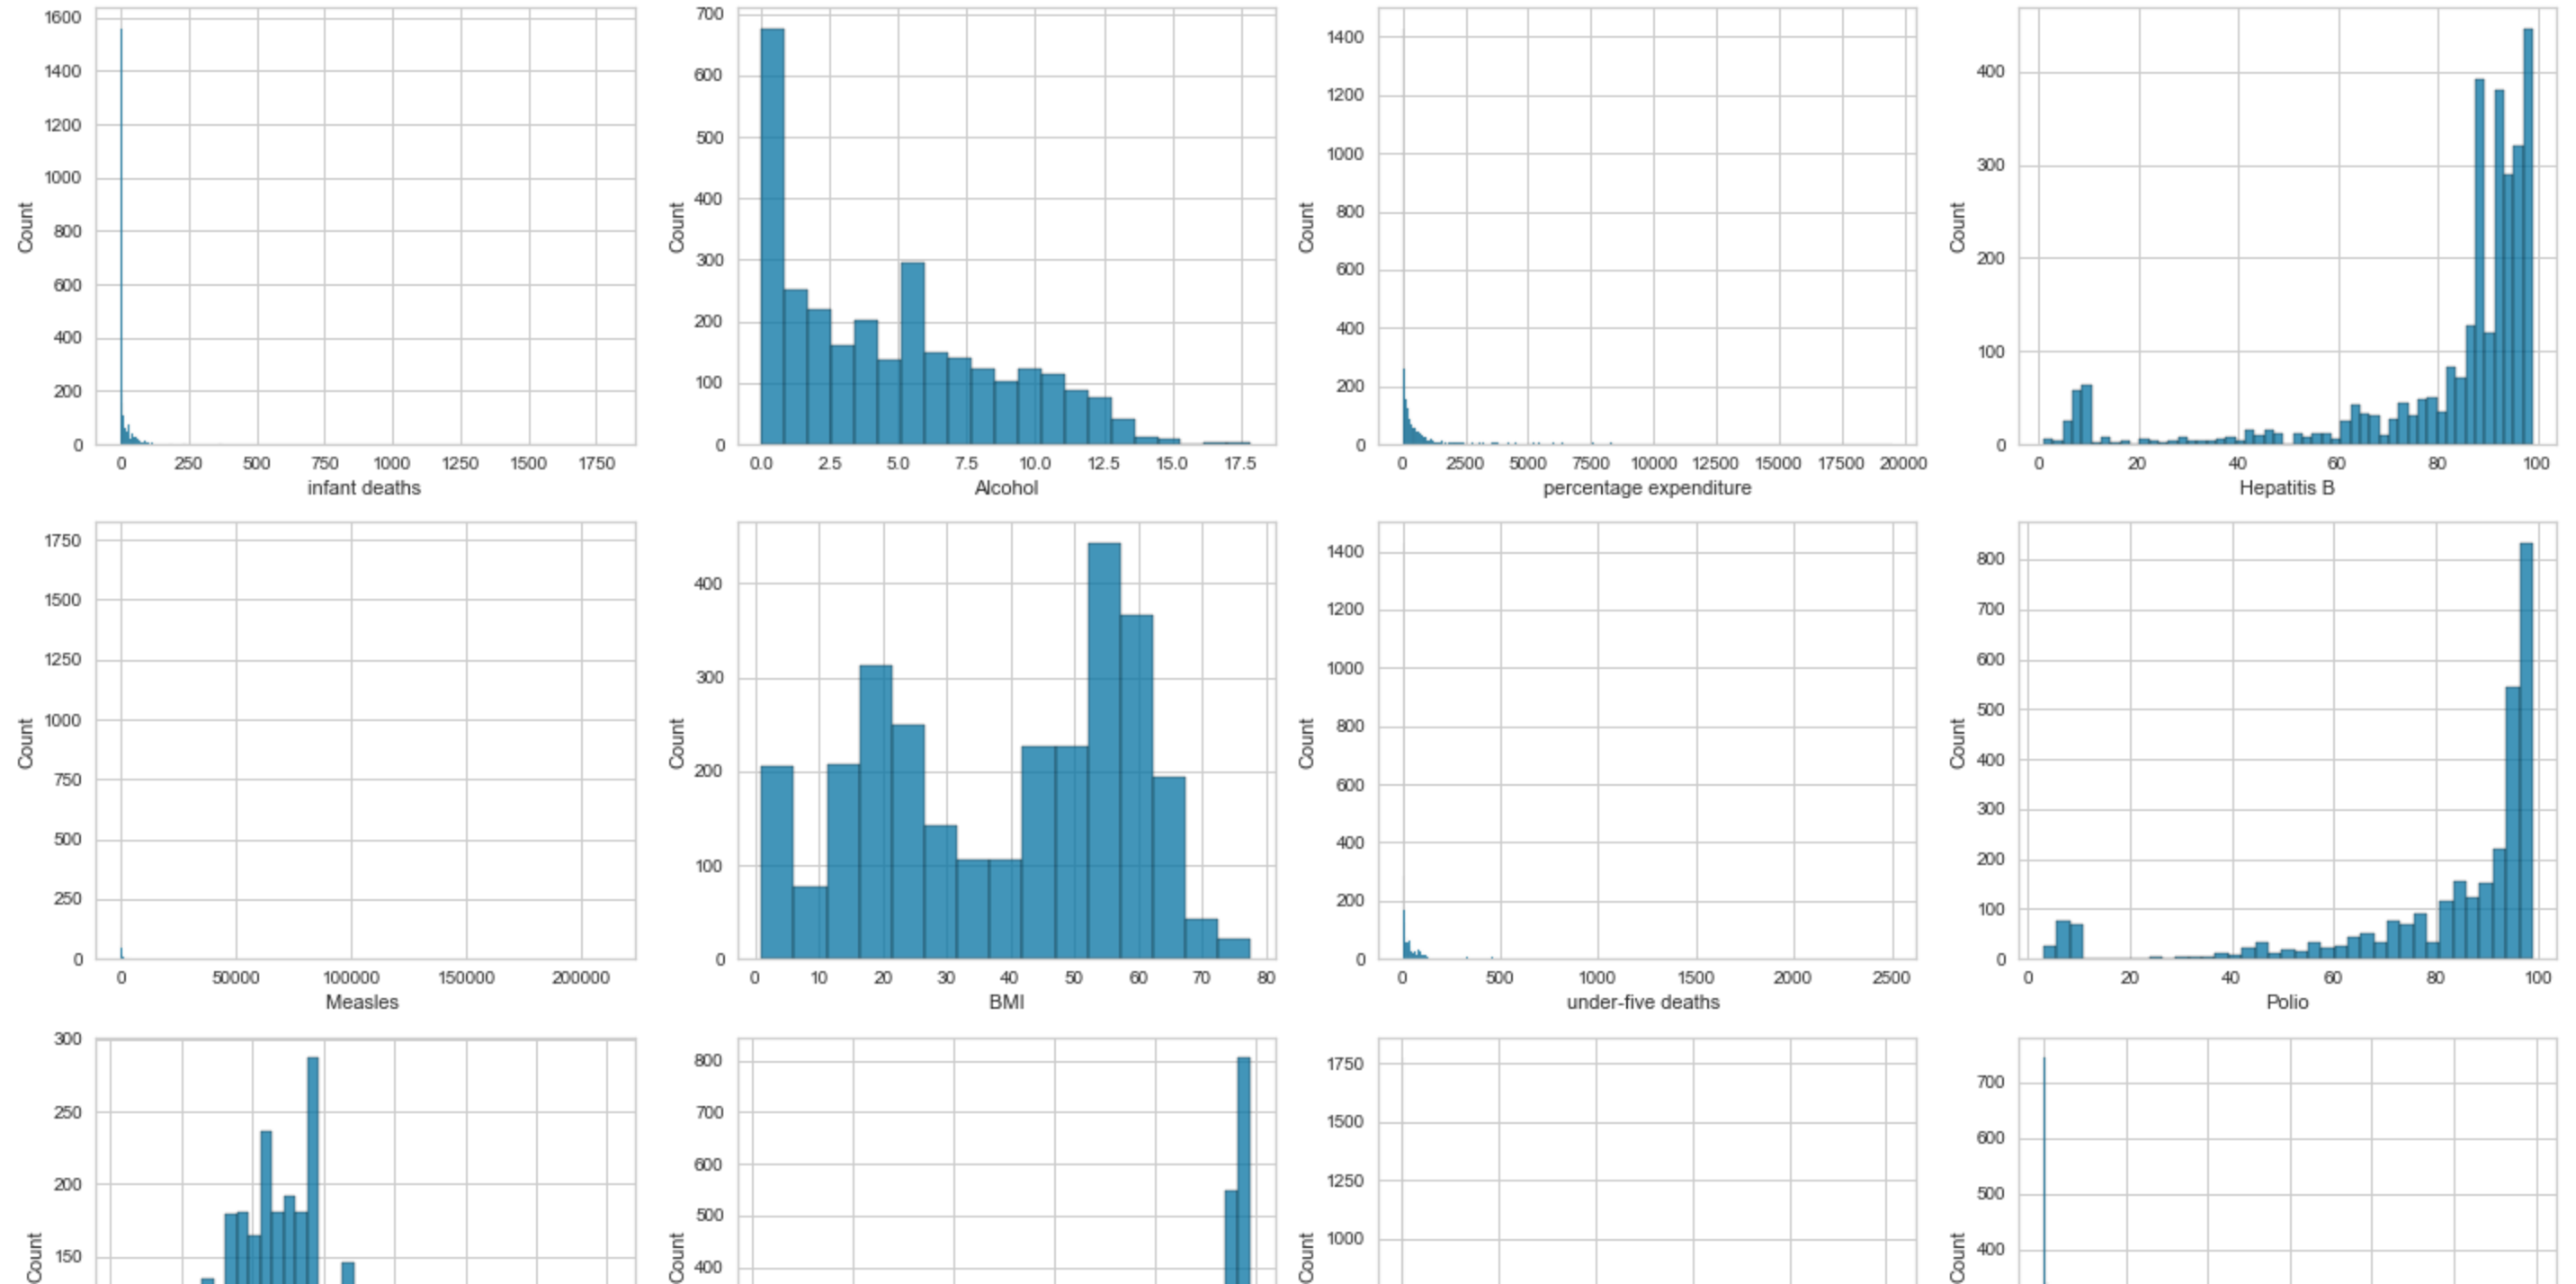
The number of NaN values in the df: 0

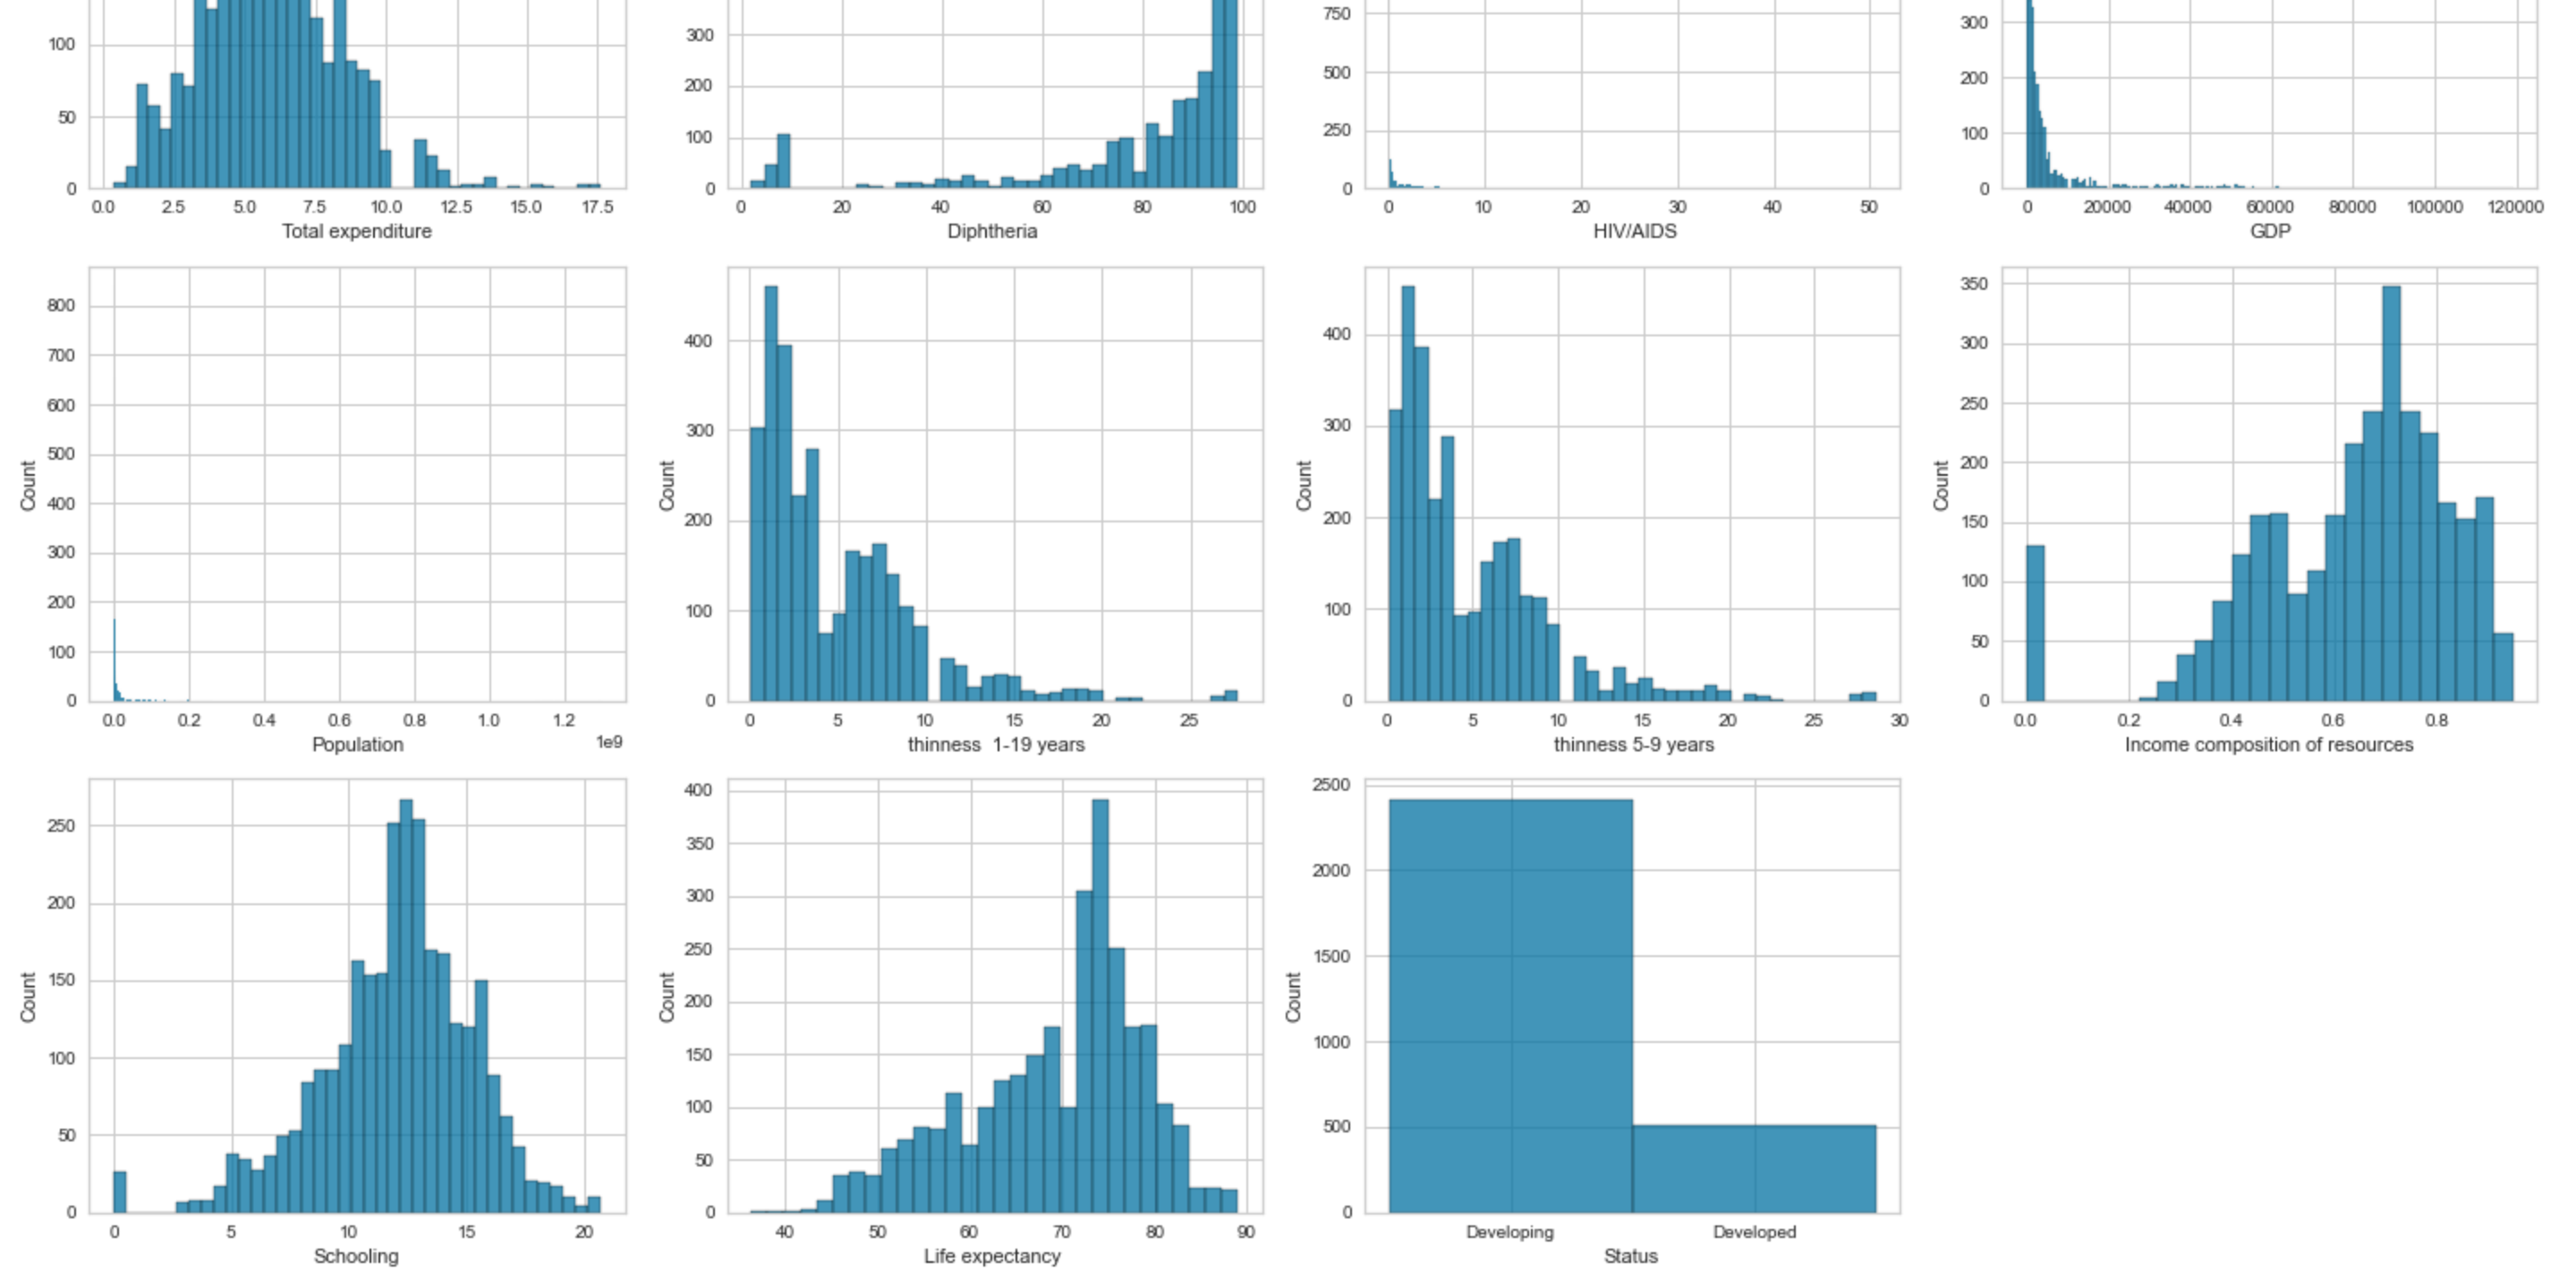
EDA:

1. Basic Plots

1.1) Counting Plots


```
In [15]: 1 col_list = ['Year', 'Adult Mortality', 'infant deaths', 'Alcohol', 'percentage expenditure', 'Hepatitis B', 'Measles ', ' BMI ', 'under-five deaths ', 'Polio', 'Total expenditure', 'Diphtheria',
2           ' thinness 1-19 years', ' thinness 5-9 years', 'Income composition of resources', 'Schooling', 'Life expectancy ', 'Status']
3
4 plt.figure(figsize=(20,20))
5 column_list = col_list[2:]
6 plt_num = 1
7
8 for i in column_list:
9     if plt_num<=18:
10         plt.subplot(5, 4, plt_num)
11         sns.histplot(df[i])
12         plt_num = plt_num+1
13     else:
14         plt.subplot(5, 4, plt_num)
15         sns.histplot(df[i])
16         plt_num = plt_num+1
17 plt.tight_layout()
```

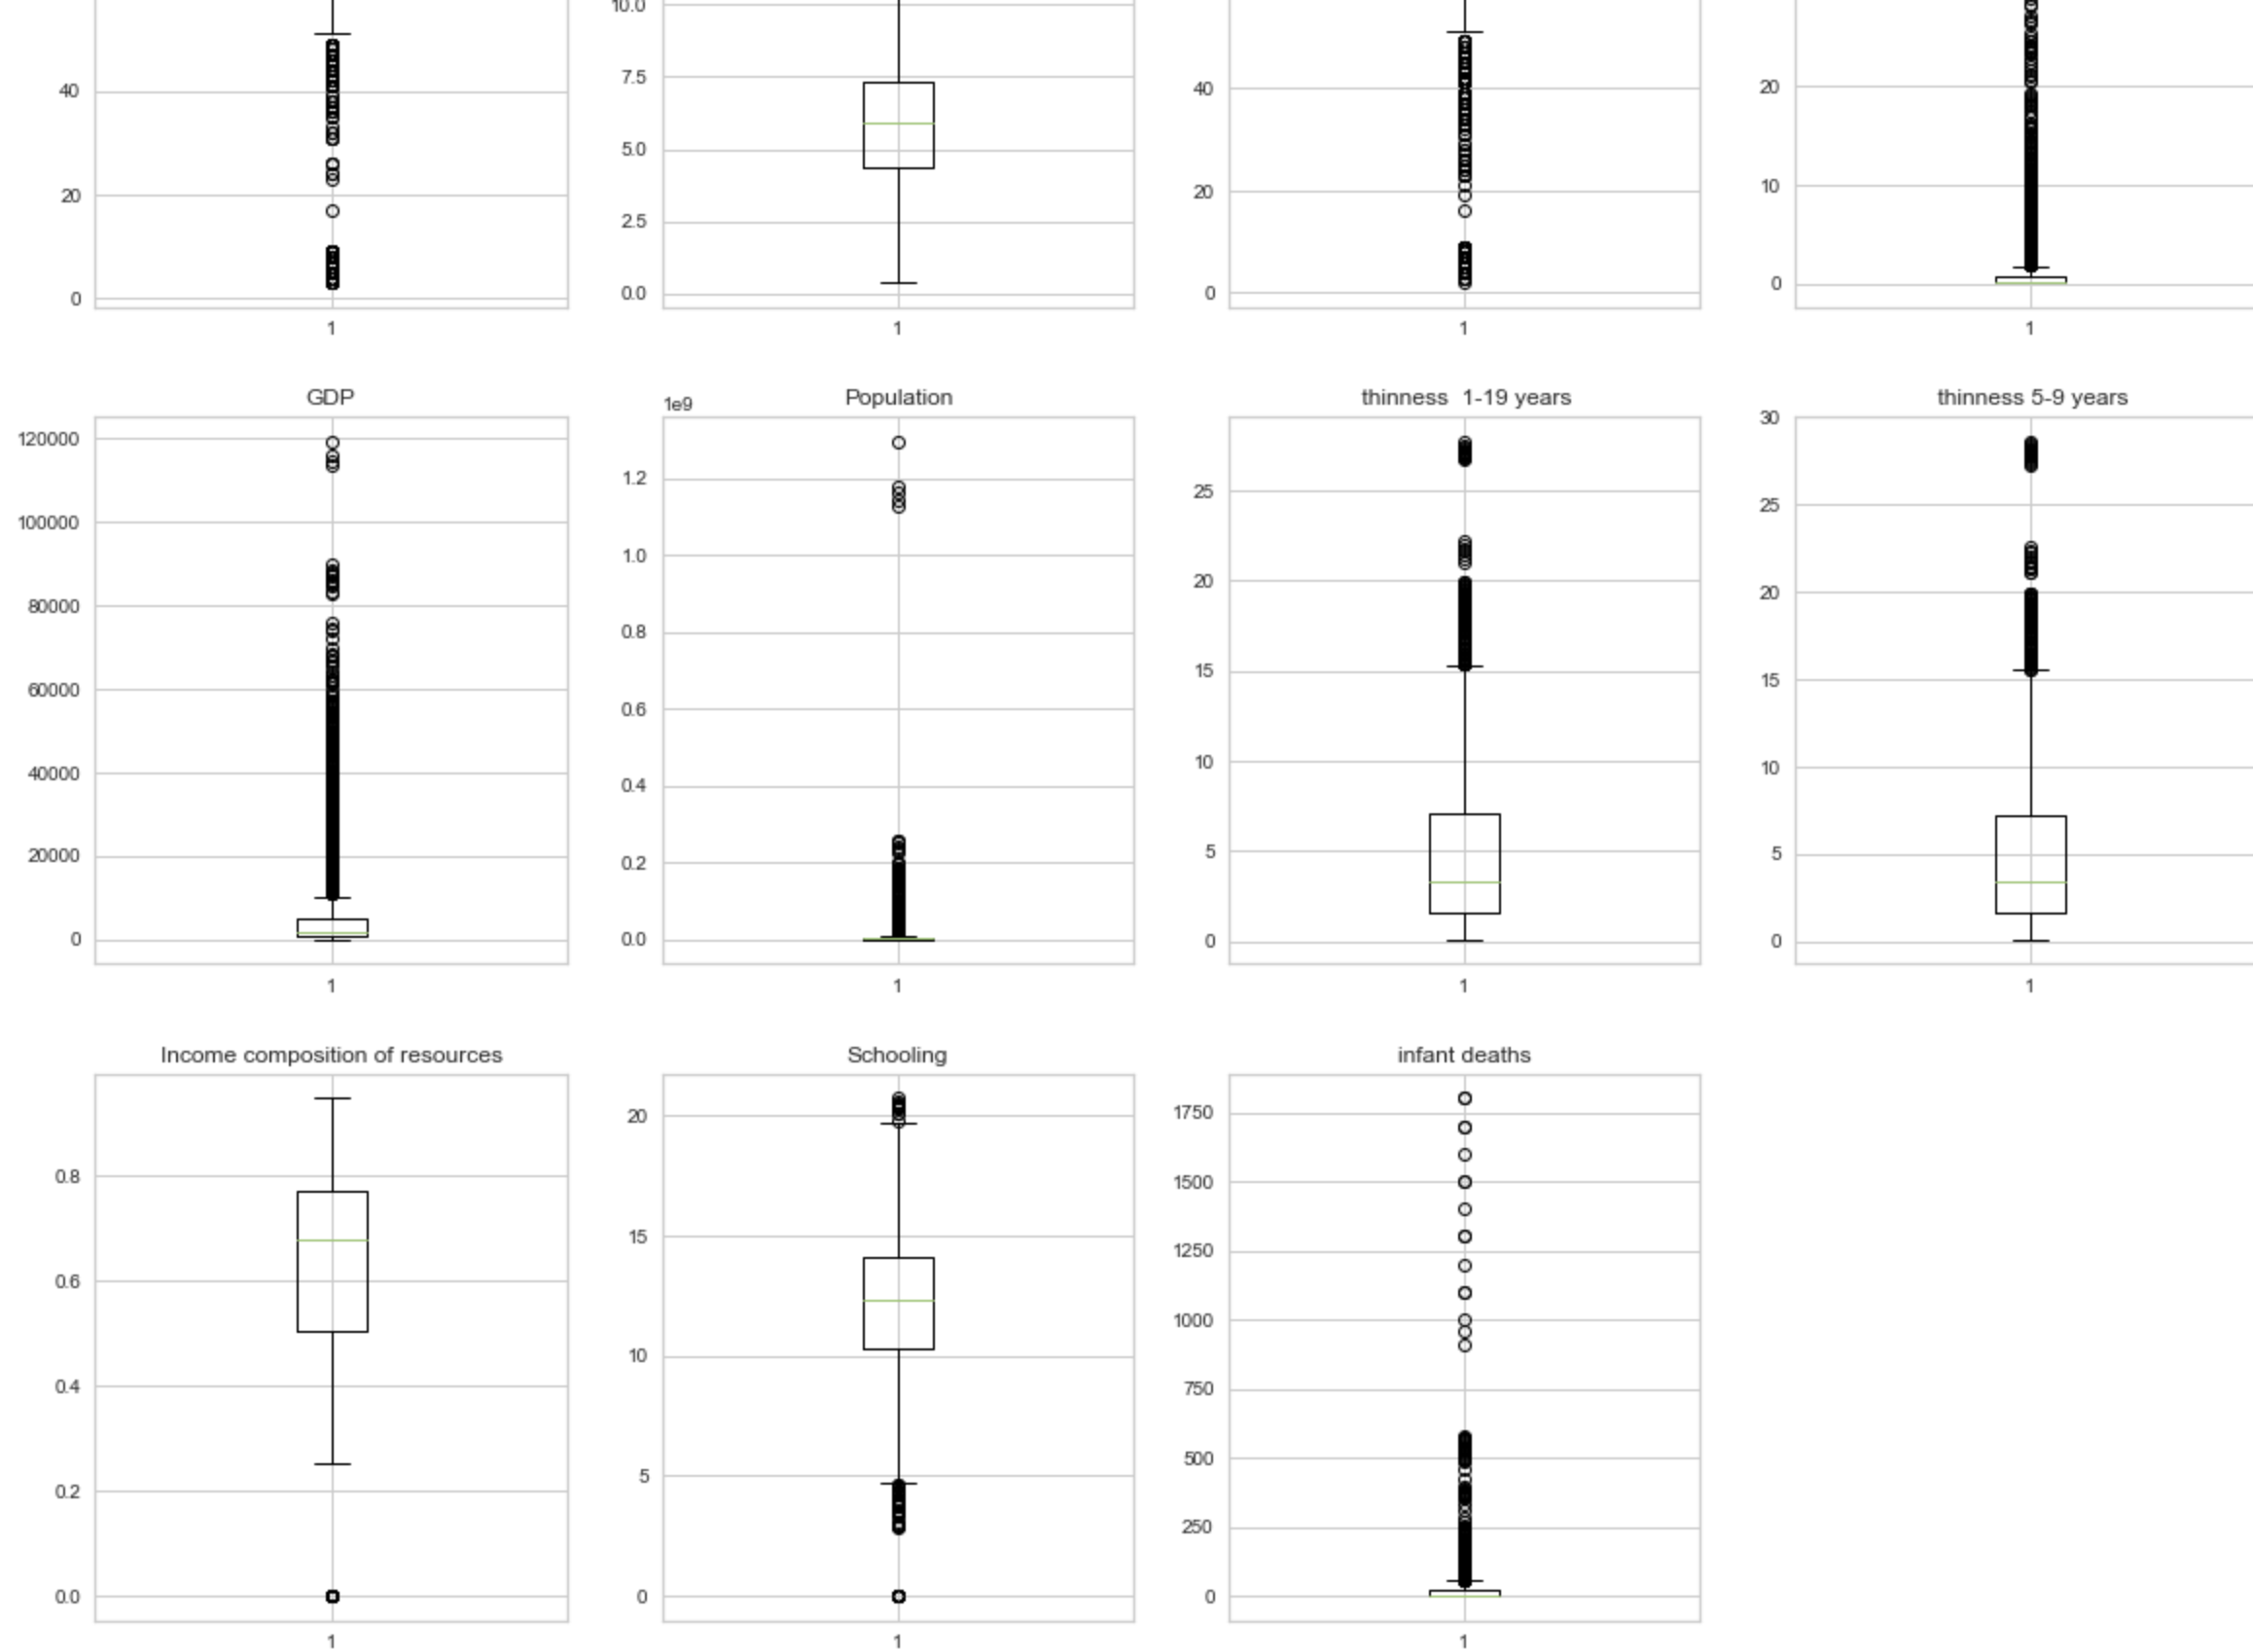




1.2) Box Plots for Outliers


```
In [16]: 1 col_dict = {'Life expectancy ': 1, 'Adult Mortality': 2, 'Alcohol': 3, 'percentage expenditure': 4, 'Hepatitis B': 5, 'Measles ': 6, ' BMI ': 7, 'under-five deaths ': 8, 'Polio' : 9, 'Total ex
2         'Diphtheria ': 11, ' HIV/AIDS': 12, 'GDP': 13, 'Population': 14, ' thinness 1-19 years': 15, ' thinness 5-9 years': 16, 'Income composition of resources': 17, 'Schooling': 18, 'infant
3
4 plt.figure(figsize=(20,30))
5
6 for variable,i in col_dict.items():
7     plt.subplot(5,4,i)
8     plt.boxplot(df[variable],whis=1.5)
9     plt.title(variable)
10
11 plt.show()
```





2. Automated EDA

```
In [17]: 1 # pip install dataprep
```

```
In [18]: 1 '''
2 from dataprep.eda import create_report
3 report = create_report(df, title='My Report')
4 report
5 '''
```

Out[18]: "\nfrom dataprep.eda import create_report\nreport = create_report(df, title='My Report')\nreport\n"

Overview

Dataset Statistics

Number of Variables	205
Number of Rows	2928
Missing Cells	0
Missing Cells (%)	0.0%
Duplicate Rows	0
Duplicate Rows (%)	0.0%
Total Size in Memory	4.6 MB
Average Row Size in Memory	1.6 KB
Variable Types	Numerical: 20 Categorical: 185

Dataset Insights

Polio and Diphtheria have similar distributions

Similar Distribution

thinness 1-19 years and thinness 5-9 years have similar distributions

Similar Distribution

infant deaths is skewed

Skewed

Alcohol is skewed

Skewed

percentage expenditure is skewed

Skewed

Hepatitis B is skewed

Skewed

Measles is skewed

Skewed

under-five deaths is skewed

Skewed

Polio is skewed

Skewed

Diphtheria is skewed


Skewed

1234567

3. Manual Plotting

3.1) Life Expectancy vs Year

```
In [19]: 1 plt.figure(figsize=(7,5))
2 plt.bar(df.groupby('Year')['Year'].count().index, df.groupby('Year')['Life expectancy '].mean(),color='red',alpha=0.65)
3 plt.xlabel("Year",fontsize=12)
4 plt.ylabel("Avg Life_Expectancy",fontsize=12)
5 plt.title("Life Expectancy w.r.t Year")
6 plt.show()
```



3.2) Life Expectancy vs Hepatitis B

The Developed countries' regression line is slightly tilted downwards stating that the Developed countries' status on getting the vaccine for Hepatitis B is gradually decreasing whereas the contrary can be noticed in the case of Developed countries.

Feature Engineering:

Label Encoding the categorical variable

In [20]:

```
1 # One Hot Encoding the Status column
2 # One Hot Encoding basically splits the unique values of a column to different columns
3
4 one_hot_encoder_status = OneHotEncoder()
5 status_resaped = np.array(df['Status']).reshape(-1, 1)
6 status_values = one_hot_encoder_status.fit_transform(status_resaped)
7
8 status_col = df['Status'].unique()
9 status_df = pd.DataFrame(status_values.toarray(), columns = status_col)
10
11 df = pd.concat([df, status_df], axis = 1)
12 df = df.drop(['Status'], axis = 1)
```

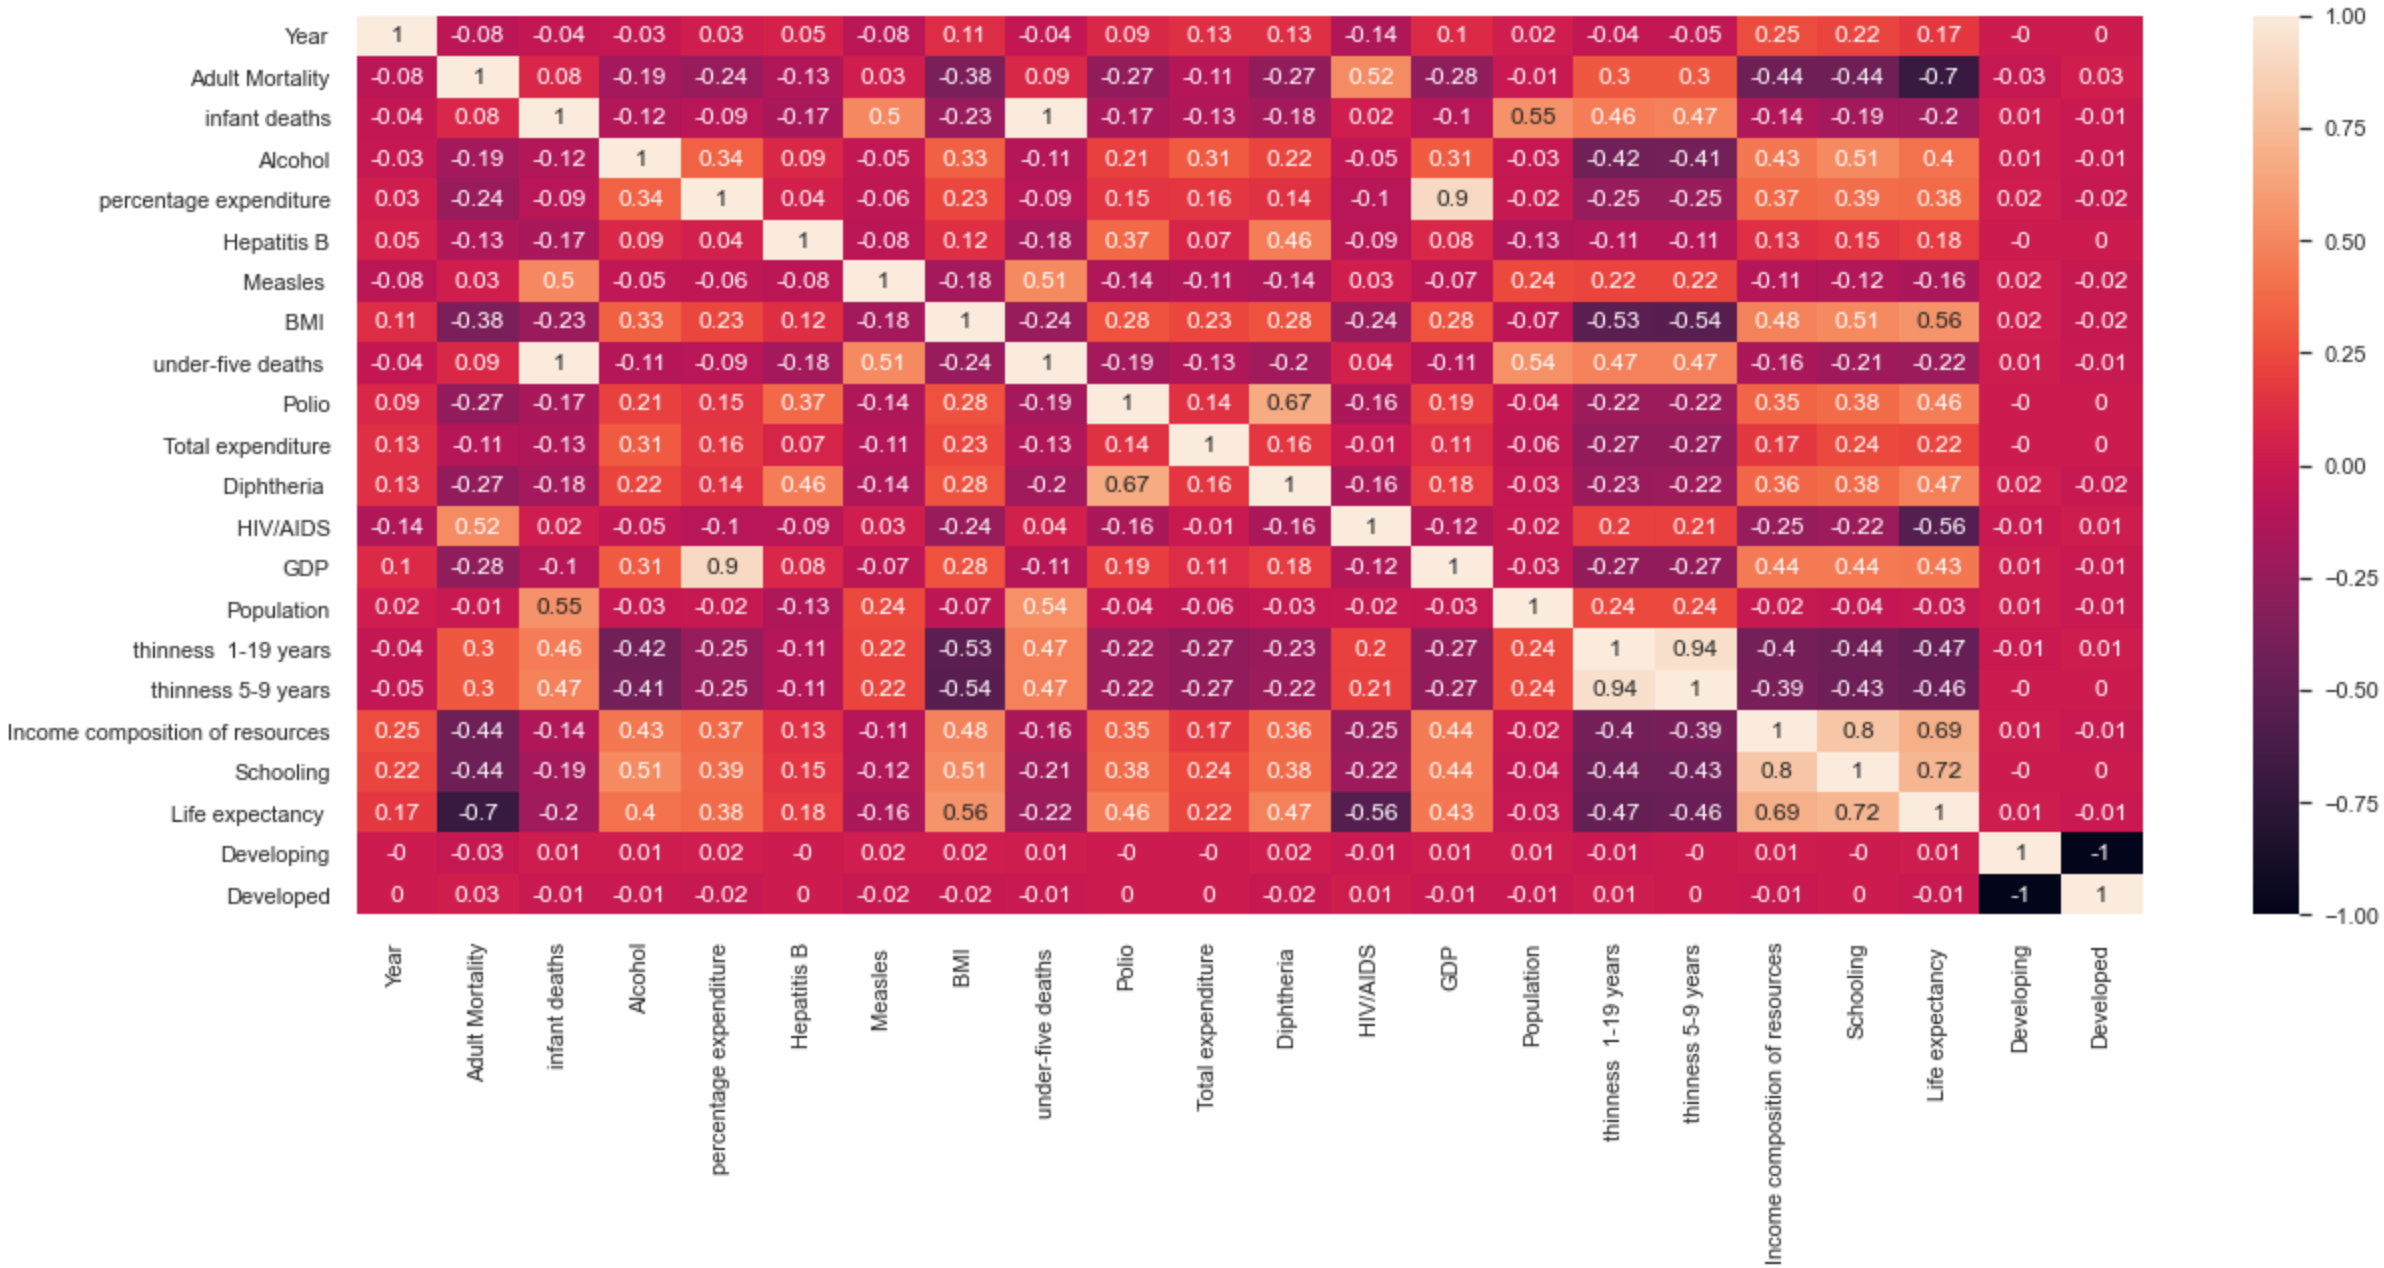
In [21]:

```
1 df.drop(['index'], axis = 1, inplace = True)
```

In [22]:

```
1 one_hot_encoder_country = OneHotEncoder()
2 country_resaped = np.array(df['Country']).reshape(-1, 1)
3 country_values = one_hot_encoder_country.fit_transform(country_resaped)
4
5 country_col = df['Country'].unique()
6 country_df = pd.DataFrame(country_values.toarray(), columns = country_col)
7
8 df = pd.concat([df, country_df], axis = 1)
9 df = df.drop(['Country'], axis = 1)
```

```
In [23]: 1 # Selecting columns with moderate to high correlation values
2 sns.set(rc={'figure.figsize':(20, 8.27)})
3 corr_map = sns.heatmap(df[['Year', 'Adult Mortality', 'infant deaths', 'Alcohol', 'percentage expenditure', 'Hepatitis B', 'Measles ', ' BMI ', 'under-five deaths ', 'Polio', 'Total expenditure',
4   ' HIV/AIDS', 'GDP', 'Population', ' thinness 1-19 years', ' thinness 5-9 years', 'Income composition of resources', 'Schooling', 'Life expectancy ', 'Developing', 'Developed']].corr()).
```



Outlier Analysis:

```
In [24]: 1 def outlierAnalysis(df, col_name):
2
3     df_score = pd.DataFrame(zscore(df[col_name]))
4
5     score_index = []
6     for i in range(len(df_score)):
7         if (df[col_name][i]>3) or (df[col_name][i]<-3):
8             score_index.append(i)
9         else:
10            pass
11
12     return score_index
```

```
In [25]: 1 # outlierAnalysis(df, 'Polio')
```

Each column seems to have great huge amount of rows with absolute value so, we avoid using this methd to inrease the efficiency of the model.

Training a Linear Regression model

```
In [145]: 1 # Splitting the data into Features and Targets
2
3 X = df.drop('Life expectancy ', axis = 1) # X --> Features
4 X = X.apply(pd.to_numeric, errors='coerce')
5 y = df['Life expectancy '] # y --> Target Variable
6 y = y.apply(pd.to_numeric, errors='coerce')
```

```
In [151]: 1 # Splitting the data into training data and testing data
2
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 500)
4 linear_regressor = LinearRegression()
```

```
In [152]: 1 # Fitting the Data
2 linear_regressor.fit(X_train, y_train)
3
4 # Predicting the Target Variable
5 y_pred = linear_regressor.predict(X_test)
```

Regression Plot

```
In [153]: 1 # Plotting the true and predicted value
2 fig = plt.figure(figsize =(20, 10))
3 sns.regplot(y_test, y_pred)
4 plt.xlabel("LE Test")
5 plt.ylabel("LE Predicted")
6 plt.title("True Value vs Predicted Value")
7 plt.show()
```


C:\Users\joanj\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(




```
In [154]: 1 # Predicted Valule vs Residuals
2 visualizer = ResidualsPlot(estimator = linear_regressor)
3 visualizer.fit(X, y)
4 visualizer.poof()
```

C:\Users\joanj\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(



Out[154]: <AxesSubplot:title={'center': 'Residuals for LinearRegression Model'}, xlabel='Predicted Value', ylabel='Residuals'>

Accuracy Count

```
In [155]: 1 # Chcking the accuracy of the model
2 print("The accuracy of the model is", linear_regressor.score(X_test, y_test))
```

The accuracy of the model is 0.8116355298655129

```
In [156]: 1 print("The r2 score of the model is", r2_score(y_test, y_pred))
```

The r2 score of the model is 0.8116355298655129

Increasing the accuracy

In [157]:

```
1 model = DecisionTreeRegressor()  
2  
3 regr_trans = TransformedTargetRegressor(regressor = model, transformer=QuantileTransformer(output_distribution='normal'))  
4 regr_trans.fit(X_train, y_train)  
5 yhat = regr_trans.predict(X_test)
```

In [160]:

```
1 print("The accuracy of the model is", regr_trans.score(X_test, y_test))
```

The accuracy of the model is 0.9073821902440397

In [161]:

```
1 print("The r2 score of the model is", r2_score(y_test, yhat))
```

The r2 score of the model is 0.9073821902440397

Evaluation Metrics

In [162]:

```
1 # Evaluation Metrics  
2 print("The Absolute Mean Error of the model is", mean_absolute_error(y_test, yhat))
```

The Absolute Mean Error of the model is 1.7109215017064847

In [163]:

```
1 print("The Mean Squared Error of the model is", mean_squared_error(y_test, yhat))
```

The Mean Squared Error of the model is 8.31316268486917

In [164]:

```
1 print("The R Mean Squared Error of the model is", math.sqrt(mean_squared_error(y_test, yhat)))
```

The R Mean Squared Error of the model is 2.8832555705086516

In [165]:

```
1 # Intercept of the model  
2 intercept = linear_regressor.intercept_  
3  
4 # Coefficeints of the model  
5 m = linear_regressor.coef_
```

```
In [166]: 1 print("The intercept of the model is", intercept, "\n\nThe weights of the features are as follows:\n", m)
```

The intercept of the model is 199673.12053704684

The weights of the features are as follows:

[-3.20715812e-02	-2.04798014e-02	9.15431201e-02	9.32558877e-02
1.60775542e-05	-1.18534786e-02	-2.90322988e-05	4.31807766e-02
-6.85738959e-02	2.23487580e-02	7.68092770e-02	4.30157311e-02
-4.75479693e-01	4.86038167e-05	7.33769351e-10	9.54822995e-03
-9.16181345e-02	5.93943805e+00	6.94182147e-01	-3.18380449e+05
-2.01723307e+05	2.16992830e+03	2.16958476e+03	2.16887124e+03
2.16977590e+03	2.16914807e+03	2.16888231e+03	2.16876489e+03
1.18826343e+05	1.18825849e+05	2.16909364e+03	2.17020312e+03
2.17001040e+03	2.16851799e+03	2.17065351e+03	2.17004570e+03
1.18827833e+05	2.17059922e+03	2.16989482e+03	2.17047721e+03
2.16972331e+03	2.17016721e+03	2.17124476e+03	2.17047939e+03
2.16885901e+03	1.18826091e+05	2.16869723e+03	2.16956072e+03
2.16797177e+03	2.16915735e+03	2.16880001e+03	2.16848294e+03
2.16884116e+03	2.17019189e+03	2.16950918e+03	2.16947285e+03
2.16642359e+03	2.16900722e+03	2.17063617e+03	2.17021969e+03
1.18827401e+05	2.16971036e+03	1.18826214e+05	1.18825233e+05
2.16898926e+03	2.16844631e+03	2.17115949e+03	1.18827444e+05
2.16963726e+03	2.16818376e+03	2.16802926e+03	2.16681671e+03
2.16727261e+03	2.16748332e+03	2.16868263e+03	2.16826762e+03
2.16745452e+03	2.16800432e+03	2.16769945e+03	2.16844795e+03
2.16910110e+03	2.16914631e+03	2.16855869e+03	1.18824990e+05
2.16856201e+03	2.16835213e+03	2.17200027e+03	2.16991594e+03
2.16809170e+03	2.16762418e+03	2.16817679e+03	2.16700281e+03
2.16717697e+03	1.18824545e+05	1.18827150e+05	2.16945869e+03
2.16846623e+03	2.17066667e+03	2.16930260e+03	1.18826164e+05
2.17043461e+03	1.18825518e+05	2.17001207e+03	1.18828080e+05
2.16919594e+03	2.16852756e+03	2.16922777e+03	2.16971917e+03
2.16938384e+03	2.16925451e+03	2.16857990e+03	1.18826921e+05
2.16999131e+03	2.16787513e+03	2.16899130e+03	2.17068234e+03
1.18825497e+05	1.18826964e+05	2.16806527e+03	2.16867295e+03
2.17008695e+03	2.16803738e+03	2.16708830e+03	1.18825162e+05
2.16658882e+03	2.16627979e+03	2.16846368e+03	2.16781526e+03
2.16702554e+03	2.16946289e+03	2.16954914e+03	2.17109144e+03
2.16961113e+03	2.17009186e+03	2.16990630e+03	1.18827589e+05
1.18827266e+05	2.16822930e+03	2.16935024e+03	2.16912302e+03
1.18825851e+05	2.16956416e+03	2.17022601e+03	2.17009413e+03
2.17114685e+03	2.16918728e+03	2.16906738e+03	2.16879851e+03
1.18825129e+05	1.18826684e+05	2.16824461e+03	2.16787666e+03
2.16986038e+03	1.18827263e+05	2.16837901e+03	2.17044041e+03
2.16963167e+03	2.16906490e+03	2.17075757e+03	2.17110481e+03
2.17119573e+03	2.16979956e+03	2.16948012e+03	2.17035024e+03
2.16960022e+03	1.18826870e+05	1.18825401e+05	1.18826866e+05
2.17127274e+03	2.16996059e+03	2.16998505e+03	2.17028118e+03
1.18825159e+05	2.16895375e+03	2.16967438e+03	2.17092925e+03
2.16980103e+03	1.18826490e+05	1.18826944e+05	2.16882894e+03
2.17095943e+03	2.16962722e+03	2.16802887e+03	2.16841266e+03
2.16798328e+03	2.16789270e+03	2.16699305e+03	2.16863527e+03
2.17017040e+03	2.16974843e+03	2.16876126e+03	2.16804349e+03
2.16822345e+03	1.18824208e+05	2.16658182e+03	1.18824976e+05
2.16703640e+03	2.16816554e+03	2.16696932e+03	2.16902185e+03
2.17029609e+03	2.16752145e+03	2.17080398e+03	2.17020757e+03]

Equation of the model

In [167]:

```
1 # Equation of the model
2
3 def predictLE(weights, x):
4
5     equation = 0
6     for i in range(len(m)):
7         equation += (weights[i]*x[i])
8     equation += intercept
9     # print("The Life Expectancy is " + str(equation))
10    return equation
```

Training a Ridge Regression model

In [168]:

```
1 # Fitting the Data
2 ridge = Ridge()
3 ridge.fit(X_train, y_train)
4
5 # Predicting the Target Variable
6 y_pred2 = ridge.predict(X_test)
```

C:\Users\joanj\anaconda3\lib\site-packages\sklearn\linear_model_ridge.py:212: LinAlgWarning: Ill-conditioned matrix (rcond=1.48442e-19): result may not be accurate.
return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T

Regression Plot

In [169]:

```
1 # Plotting the true and predicted value
2 fig = plt.figure(figsize =(20, 10))
3 sns.regplot(y_test, y_pred2)
4 plt.xlabel("LE Test")
5 plt.ylabel("LE Predicted")
6 plt.title("True Value vs Predicted Value")
7 plt.show()
```

C:\Users\joanj\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



Accuracy Count

```
In [170]: 1 # Chcking the accuracy of the model
          2 print("The accuracy of the model is", ridge.score(X_test, y_test))
```

The accuracy of the model is 0.8139971860712808

```
In [171]: 1 print("The r2 score of the model is", r2_score(y_test, y_pred2))
```

The r2 score of the model is 0.8139971860712808

Evaluation Metrics

```
In [172]: 1 # Evaluation Metrics
          2 print("The Absolute Mean Error of the model is", mean_absolute_error(y_test, y_pred2))
```

The Absolute Mean Error of the model is 3.0948610206583105

```
In [173]: 1 print("The Mean Squared Error of the model is", mean_squared_error(y_test, y_pred2))
```

The Mean Squared Error of the model is 16.695186985172498

```
In [174]: 1 print("The RMSE of the model is", math.sqrt(mean_squared_error(y_test, y_pred2)))
```

The RMSE of the model is 4.0859744229709145

```
In [175]: 1 # Intercept of the model
          2 intercept2 = linear_regressor.intercept_
          3
          4 # Coefficeints of the model
          5 m2 = linear_regressor.coef_
```

```
In [176]: 1 print("The intercept of the model is", intercept2, "\n\nThe weights of the features are as follows:\n", m2)
```

The intercept of the model is 199673.12053704684

The weights of the features are as follows:

```
[-3.20715812e-02 -2.04798014e-02  9.15431201e-02  9.32558877e-02
 1.60775542e-05 -1.18534786e-02 -2.90322988e-05  4.31807766e-02
-6.85738959e-02  2.23487580e-02  7.68092770e-02  4.30157311e-02
-4.75479693e-01  4.86038167e-05  7.33769351e-10  9.54822995e-03
-9.16181345e-02  5.93943805e+00  6.94182147e-01 -3.18380449e+05
-2.01723307e+05  2.16992830e+03  2.16958476e+03  2.16887124e+03
 2.16977590e+03  2.16914807e+03  2.16888231e+03  2.16876489e+03
 1.18826343e+05  1.18825849e+05  2.16909364e+03  2.17020312e+03
 2.17001040e+03  2.16851799e+03  2.17065351e+03  2.17004570e+03
 1.18827833e+05  2.17059922e+03  2.16989482e+03  2.17047721e+03
 2.16972331e+03  2.17016721e+03  2.17124476e+03  2.17047939e+03
 2.16885901e+03  1.18826091e+05  2.16869723e+03  2.16956072e+03
 2.16797177e+03  2.16915735e+03  2.16880001e+03  2.16848294e+03
 2.16884116e+03  2.17019189e+03  2.16950918e+03  2.16947285e+03
 2.16642359e+03  2.16900722e+03  2.17063617e+03  2.17021969e+03
 1.18827401e+05  2.16971036e+03  1.18826214e+05  1.18825233e+05
 2.16898926e+03  2.16844631e+03  2.17115949e+03  1.18827444e+05
 2.16963726e+03  2.16818376e+03  2.16802926e+03  2.16681671e+03
 2.16727261e+03  2.16748332e+03  2.16868263e+03  2.16826762e+03
 2.16745452e+03  2.16800432e+03  2.16769945e+03  2.16844795e+03
 2.16910110e+03  2.16914631e+03  2.16855869e+03  1.18824990e+05
 2.16856201e+03  2.16835213e+03  2.17200027e+03  2.16991594e+03
 2.16809170e+03  2.16762418e+03  2.16817679e+03  2.16700281e+03
 2.16717697e+03  1.18824545e+05  1.18827150e+05  2.16945869e+03
 2.16846623e+03  2.17066667e+03  2.16930260e+03  1.18826164e+05
 2.17043461e+03  1.18825518e+05  2.17001207e+03  1.18828080e+05
 2.16919594e+03  2.16852756e+03  2.16922777e+03  2.16971917e+03
 2.16938384e+03  2.16925451e+03  2.16857990e+03  1.18826921e+05
 2.16999131e+03  2.16787513e+03  2.16899130e+03  2.17068234e+03
 1.18825497e+05  1.18826964e+05  2.16806527e+03  2.16867295e+03
 2.17008695e+03  2.16803738e+03  2.16708830e+03  1.18825162e+05
 2.16658882e+03  2.16627979e+03  2.16846368e+03  2.16781526e+03
 2.16702554e+03  2.16946289e+03  2.16954914e+03  2.17109144e+03
 2.16961113e+03  2.17009186e+03  2.16990630e+03  1.18827589e+05
 1.18827266e+05  2.16822930e+03  2.16935024e+03  2.16912302e+03
 1.18825851e+05  2.16956416e+03  2.17022601e+03  2.17009413e+03
 2.17114685e+03  2.16918728e+03  2.16906738e+03  2.16879851e+03
 1.18825129e+05  1.18826684e+05  2.16824461e+03  2.16787666e+03
 2.16986038e+03  1.18827263e+05  2.16837901e+03  2.17044041e+03
 2.16963167e+03  2.16906490e+03  2.17075757e+03  2.17110481e+03
 2.17119573e+03  2.16979956e+03  2.16948012e+03  2.17035024e+03
 2.16960022e+03  1.18826870e+05  1.18825401e+05  1.18826866e+05
 2.17127274e+03  2.16996059e+03  2.16998505e+03  2.17028118e+03
 1.18825159e+05  2.16895375e+03  2.16967438e+03  2.17092925e+03
 2.16980103e+03  1.18826490e+05  1.18826944e+05  2.16882894e+03
 2.17095943e+03  2.16962722e+03  2.16802887e+03  2.16841266e+03
 2.16798328e+03  2.16789270e+03  2.16699305e+03  2.16863527e+03
 2.17017040e+03  2.16974843e+03  2.16876126e+03  2.16804349e+03
 2.16822345e+03  1.18824208e+05  2.16658182e+03  1.18824976e+05
 2.16703640e+03  2.16816554e+03  2.16696932e+03  2.16902185e+03
 2.17029609e+03  2.16752145e+03  2.17080398e+03  2.17020757e+03]
```

```
In [177]: 1 # Equation of the model
2
3 def predictLE2(weights, x):
4
5     equation = 0
6     for i in range(len(m2)):
7         equation += (weights[i]*x[i])
8     equation += intercept2
9     print("The Life Expectancy is " + str(equation))
```

Training a Decision Tree model


```
In [178]: 1 # Fitting the Data
2 tree = DecisionTreeRegressor(max_depth=4)
3 tree.fit(X_train, y_train)
4
5 # Predicting the Target Variable
6 y_pred3 = tree.predict(X_test)
```

Regression Plot


```
In [179]: 1 # Plotting the true and predicted value
2 fig = plt.figure(figsize =(20, 10))
3 sns.regplot(y_test, y_pred3)
4 plt.xlabel("LE Test")
5 plt.ylabel("LE Predicted")
6 plt.title("True Value vs Predicted Value")
7 plt.show()
```


C:\Users\joanj\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

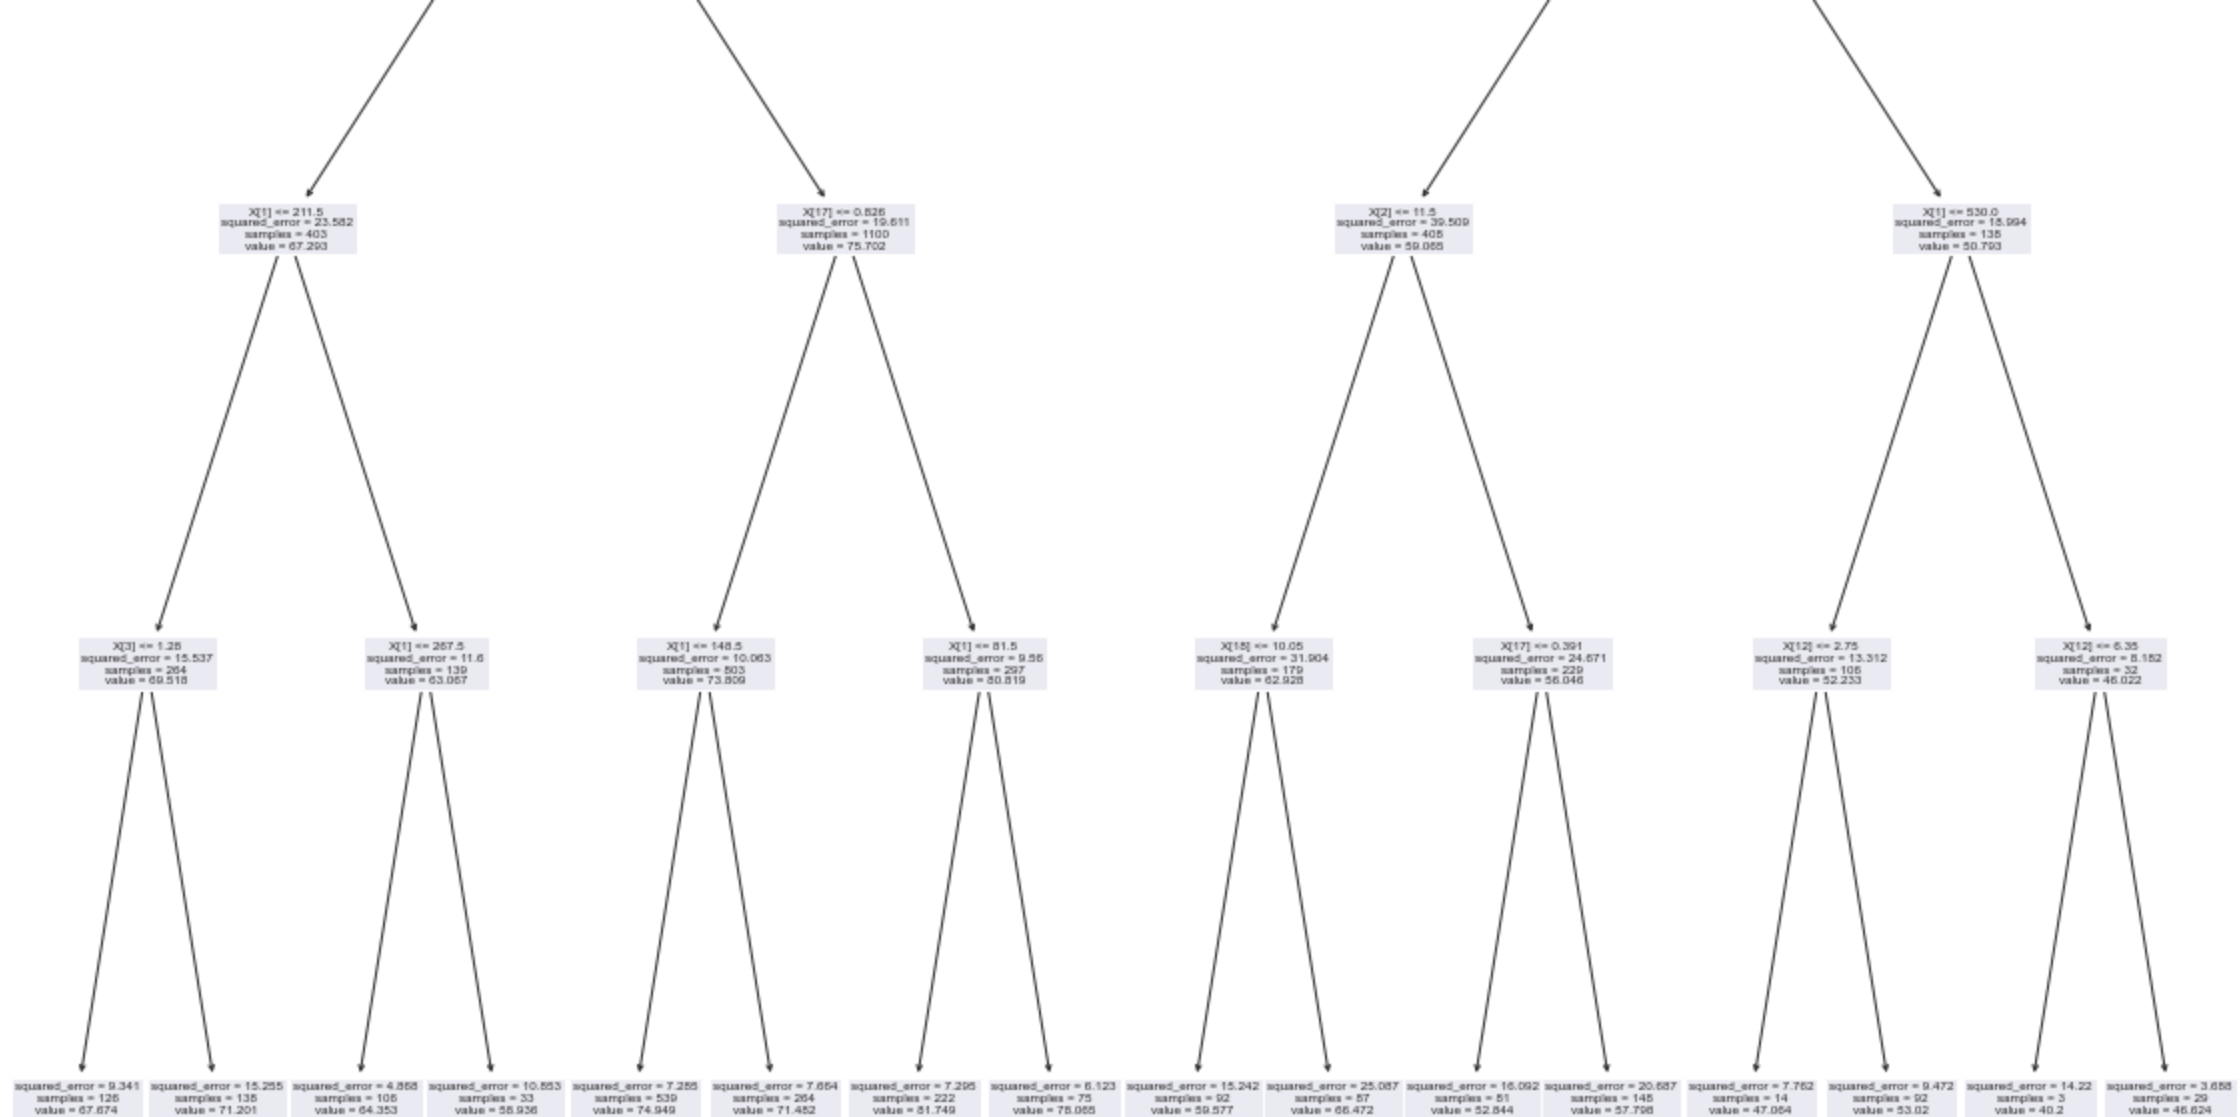
```
warnings.warn(
```




```
In [180]: 1 # Show the decision tree
2 plt.figure(figsize=(20,20))
3 plot_tree(tree)
```

Out[180]: [Text(0.5, 0.9, 'X[12] <= 0.65\nsquared_error = 90.972\nsamples = 2049\nvalue = 69.058'),
Text(0.25, 0.7, 'X[17] <= 0.638\nsquared_error = 34.551\nsamples = 1503\nvalue = 73.447'),
Text(0.125, 0.5, 'X[1] <= 211.5\nsquared_error = 23.582\nsamples = 403\nvalue = 67.293'),
Text(0.0625, 0.3, 'X[3] <= 1.28\nsquared_error = 15.537\nsamples = 264\nvalue = 69.518'),
Text(0.03125, 0.1, 'squared_error = 9.341\nsamples = 126\nvalue = 67.674'),
Text(0.09375, 0.1, 'squared_error = 15.255\nsamples = 138\nvalue = 71.201'),
Text(0.1875, 0.3, 'X[1] <= 267.5\nsquared_error = 11.6\nsamples = 139\nvalue = 63.067'),
Text(0.15625, 0.1, 'squared_error = 4.868\nsamples = 106\nvalue = 64.353'),
Text(0.21875, 0.1, 'squared_error = 10.853\nsamples = 33\nvalue = 58.936'),
Text(0.375, 0.5, 'X[17] <= 0.826\nsquared_error = 19.611\nsamples = 1100\nvalue = 75.702'),
Text(0.3125, 0.3, 'X[1] <= 148.5\nsquared_error = 10.063\nsamples = 803\nvalue = 73.809'),
Text(0.28125, 0.1, 'squared_error = 7.285\nsamples = 539\nvalue = 74.949'),
Text(0.34375, 0.1, 'squared_error = 7.664\nsamples = 264\nvalue = 71.482'),
Text(0.4375, 0.3, 'X[1] <= 81.5\nsquared_error = 9.56\nsamples = 297\nvalue = 80.819'),
Text(0.40625, 0.1, 'squared_error = 7.295\nsamples = 222\nvalue = 81.749'),
Text(0.46875, 0.1, 'squared_error = 6.123\nsamples = 75\nvalue = 78.065'),
Text(0.75, 0.7, 'X[1] <= 374.0\nsquared_error = 47.249\nsamples = 546\nvalue = 56.975'),
Text(0.625, 0.5, 'X[2] <= 11.5\nsquared_error = 39.509\nsamples = 408\nvalue = 59.065'),
Text(0.5625, 0.3, 'X[18] <= 10.05\nsquared_error = 31.904\nsamples = 179\nvalue = 62.928'),
Text(0.53125, 0.1, 'squared_error = 15.242\nsamples = 92\nvalue = 59.577'),
Text(0.59375, 0.1, 'squared_error = 25.087\nsamples = 87\nvalue = 66.472'),
Text(0.6875, 0.3, 'X[17] <= 0.391\nsquared_error = 24.671\nsamples = 229\nvalue = 56.046'),
Text(0.65625, 0.1, 'squared_error = 16.092\nsamples = 81\nvalue = 52.844'),
Text(0.71875, 0.1, 'squared_error = 20.687\nsamples = 148\nvalue = 57.798'),
Text(0.875, 0.5, 'X[1] <= 530.0\nsquared_error = 18.994\nsamples = 138\nvalue = 50.793'),
Text(0.8125, 0.3, 'X[12] <= 2.75\nsquared_error = 13.312\nsamples = 106\nvalue = 52.233'),
Text(0.78125, 0.1, 'squared_error = 7.762\nsamples = 14\nvalue = 47.064'),
Text(0.84375, 0.1, 'squared_error = 9.472\nsamples = 92\nvalue = 53.02'),
Text(0.9375, 0.3, 'X[12] <= 6.35\nsquared_error = 8.182\nsamples = 32\nvalue = 46.022'),
Text(0.90625, 0.1, 'squared_error = 14.22\nsamples = 3\nvalue = 40.2'),
Text(0.96875, 0.1, 'squared_error = 3.688\nsamples = 29\nvalue = 46.624')]





Accuracy Count

```
In [181]: 1 # Chcking the accuracy of the model
          2 print("The accuracy of the model is", tree.score(X_test, y_test))
```

The accuracy of the model is 0.8569892944426825

```
In [182]: 1 print("The r2 score of the model is", r2_score(y_test, y_pred3))
```

The r2 score of the model is 0.8569892944426825

Evaluation Metrics

```
In [183]: 1 # Evaluation Metrics
          2 print("The Absolute Mean Error of the model is", mean_absolute_error(y_test, y_pred3))
```

The Absolute Mean Error of the model is 2.693555539173137

In [184]:

1 print("The Mean Squared Error of the model is", mean_squared_error(y_test, y_pred3))

The Mean Squared Error of the model is 12.836313708005761

In [185]:

1 print("The RMSE of the model is", math.sqrt(mean_squared_error(y_test, y_pred3)))

The RMSE of the model is 3.58278016462157

In [186]:

1 *# Intercept of the model*
2 intercept3 = linear_regressor.intercept_
3
4 *# Coefficeints of the model*
5 m3 = linear_regressor.coef_

```
In [187]: 1 print("The intercept of the model is", intercept3, "\n\nThe weights of the features are as follows:\n", m3)
```

The intercept of the model is 199673.12053704684

The weights of the features are as follows:

[-3.20715812e-02	-2.04798014e-02	9.15431201e-02	9.32558877e-02
1.60775542e-05	-1.18534786e-02	-2.90322988e-05	4.31807766e-02
-6.85738959e-02	2.23487580e-02	7.68092770e-02	4.30157311e-02
-4.75479693e-01	4.86038167e-05	7.33769351e-10	9.54822995e-03
-9.16181345e-02	5.93943805e+00	6.94182147e-01	-3.18380449e+05
-2.01723307e+05	2.16992830e+03	2.16958476e+03	2.16887124e+03
2.16977590e+03	2.16914807e+03	2.16888231e+03	2.16876489e+03
1.18826343e+05	1.18825849e+05	2.16909364e+03	2.17020312e+03
2.17001040e+03	2.16851799e+03	2.17065351e+03	2.17004570e+03
1.18827833e+05	2.17059922e+03	2.16989482e+03	2.17047721e+03
2.16972331e+03	2.17016721e+03	2.17124476e+03	2.17047939e+03
2.16885901e+03	1.18826091e+05	2.16869723e+03	2.16956072e+03
2.16797177e+03	2.16915735e+03	2.16880001e+03	2.16848294e+03
2.16884116e+03	2.17019189e+03	2.16950918e+03	2.16947285e+03
2.16642359e+03	2.16900722e+03	2.17063617e+03	2.17021969e+03
1.18827401e+05	2.16971036e+03	1.18826214e+05	1.18825233e+05
2.16898926e+03	2.16844631e+03	2.17115949e+03	1.18827444e+05
2.16963726e+03	2.16818376e+03	2.16802926e+03	2.16681671e+03
2.16727261e+03	2.16748332e+03	2.16868263e+03	2.16826762e+03
2.16745452e+03	2.16800432e+03	2.16769945e+03	2.16844795e+03
2.16910110e+03	2.16914631e+03	2.16855869e+03	1.18824990e+05
2.16856201e+03	2.16835213e+03	2.17200027e+03	2.16991594e+03
2.16809170e+03	2.16762418e+03	2.16817679e+03	2.16700281e+03
2.16717697e+03	1.18824545e+05	1.18827150e+05	2.16945869e+03
2.16846623e+03	2.17066667e+03	2.16930260e+03	1.18826164e+05
2.17043461e+03	1.18825518e+05	2.17001207e+03	1.18828080e+05
2.16919594e+03	2.16852756e+03	2.16922777e+03	2.16971917e+03
2.16938384e+03	2.16925451e+03	2.16857990e+03	1.18826921e+05
2.16999131e+03	2.16787513e+03	2.16899130e+03	2.17068234e+03
1.18825497e+05	1.18826964e+05	2.16806527e+03	2.16867295e+03
2.17008695e+03	2.16803738e+03	2.16708830e+03	1.18825162e+05
2.16658882e+03	2.16627979e+03	2.16846368e+03	2.16781526e+03
2.16702554e+03	2.16946289e+03	2.16954914e+03	2.17109144e+03
2.16961113e+03	2.17009186e+03	2.16990630e+03	1.18827589e+05
1.18827266e+05	2.16822930e+03	2.16935024e+03	2.16912302e+03
1.18825851e+05	2.16956416e+03	2.17022601e+03	2.17009413e+03
2.17114685e+03	2.16918728e+03	2.16906738e+03	2.16879851e+03
1.18825129e+05	1.18826684e+05	2.16824461e+03	2.16787666e+03
2.16986038e+03	1.18827263e+05	2.16837901e+03	2.17044041e+03
2.16963167e+03	2.16906490e+03	2.17075757e+03	2.17110481e+03
2.17119573e+03	2.16979956e+03	2.16948012e+03	2.17035024e+03
2.16960022e+03	1.18826870e+05	1.18825401e+05	1.18826866e+05
2.17127274e+03	2.16996059e+03	2.16998505e+03	2.17028118e+03
1.18825159e+05	2.16895375e+03	2.16967438e+03	2.17092925e+03
2.16980103e+03	1.18826490e+05	1.18826944e+05	2.16882894e+03
2.17095943e+03	2.16962722e+03	2.16802887e+03	2.16841266e+03
2.16798328e+03	2.16789270e+03	2.16699305e+03	2.16863527e+03
2.17017040e+03	2.16974843e+03	2.16876126e+03	2.16804349e+03
2.16822345e+03	1.18824208e+05	2.16658182e+03	1.18824976e+05
2.16703640e+03	2.16816554e+03	2.16696932e+03	2.16902185e+03
2.17029609e+03	2.16752145e+03	2.17080398e+03	2.17020757e+03]

```
In [188]: 1 # Equation of the model
2
3 def predictLE3(weights, x):
4
5     equation = 0
6     for i in range(len(m3)):
7         equation += (weights[i]*x[i])
8     equation += intercept3
9     print("The Life Expectancy is " + str(equation))
```

Training a Random Forest model

```
In [189]: 1 # Fitting the Data
2 forest = RandomForestRegressor(n_estimators=100)
3 forest.fit(X_train, y_train)
4
5 # Predicting the Target Variable
6 y_pred4 = forest.predict(X_test)
```


Regression Plot

In [190]:

```
1 # Plotting the true and predicted value
2 fig = plt.figure(figsize =(20, 10))
3 sns.regplot(y_test, y_pred4)
4 plt.xlabel("LE Test")
5 plt.ylabel("LE Predicted")
6 plt.title("True Value vs Predicted Value")
7 plt.show()
```

C:\Users\joanj\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



Accuracy Count

In [191]:

1

Chcking the accuracy of the model

2

print("The accuracy of the model is", forest.score(X_test, y_test))

The accuracy of the model is 0.9558929929973955

In [192]:

1

print("The r2 score of the model is", r2_score(y_test, y_pred4))

The r2 score of the model is 0.9558929929973955

Evaluation Metrics

In [193]:

1

Evaluation Metrics

2

print("The Absolute Mean Error of the model is", mean_absolute_error(y_test, y_pred4))

The Absolute Mean Error of the model is 1.2947281001137665

In [194]:

1

print("The Mean Squared Error of the model is", mean_squared_error(y_test, y_pred4))

The Mean Squared Error of the model is 3.9589440273037573

In [195]:

1

print("The Mean Squared Error of the model is", math.sqrt(mean_squared_error(y_test, y_pred4)))

The Mean Squared Error of the model is 1.989709533400229

In [196]:

1

Intercept of the model

2

intercept4 = linear_regressor.intercept_

3

4

Coefficeints of the model

5

m4 = linear_regressor.coef_

```
In [197]: 1 print("The intercept of the model is", intercept4, "\n\nThe weights of the features are as follows:\n", m4)
```

The intercept of the model is 199673.12053704684

The weights of the features are as follows:

[-3.20715812e-02	-2.04798014e-02	9.15431201e-02	9.32558877e-02
1.60775542e-05	-1.18534786e-02	-2.90322988e-05	4.31807766e-02
-6.85738959e-02	2.23487580e-02	7.68092770e-02	4.30157311e-02
-4.75479693e-01	4.86038167e-05	7.33769351e-10	9.54822995e-03
-9.16181345e-02	5.93943805e+00	6.94182147e-01	-3.18380449e+05
-2.01723307e+05	2.16992830e+03	2.16958476e+03	2.16887124e+03
2.16977590e+03	2.16914807e+03	2.16888231e+03	2.16876489e+03
1.18826343e+05	1.18825849e+05	2.16909364e+03	2.17020312e+03
2.17001040e+03	2.16851799e+03	2.17065351e+03	2.17004570e+03
1.18827833e+05	2.17059922e+03	2.16989482e+03	2.17047721e+03
2.16972331e+03	2.17016721e+03	2.17124476e+03	2.17047939e+03
2.16885901e+03	1.18826091e+05	2.16869723e+03	2.16956072e+03
2.16797177e+03	2.16915735e+03	2.16880001e+03	2.16848294e+03
2.16884116e+03	2.17019189e+03	2.16950918e+03	2.16947285e+03
2.16642359e+03	2.16900722e+03	2.17063617e+03	2.17021969e+03
1.18827401e+05	2.16971036e+03	1.18826214e+05	1.18825233e+05
2.16898926e+03	2.16844631e+03	2.17115949e+03	1.18827444e+05
2.16963726e+03	2.16818376e+03	2.16802926e+03	2.16681671e+03
2.16727261e+03	2.16748332e+03	2.16868263e+03	2.16826762e+03
2.16745452e+03	2.16800432e+03	2.16769945e+03	2.16844795e+03
2.16910110e+03	2.16914631e+03	2.16855869e+03	1.18824990e+05
2.16856201e+03	2.16835213e+03	2.17200027e+03	2.16991594e+03
2.16809170e+03	2.16762418e+03	2.16817679e+03	2.16700281e+03
2.16717697e+03	1.18824545e+05	1.18827150e+05	2.16945869e+03
2.16846623e+03	2.17066667e+03	2.16930260e+03	1.18826164e+05
2.17043461e+03	1.18825518e+05	2.17001207e+03	1.18828080e+05
2.16919594e+03	2.16852756e+03	2.16922777e+03	2.16971917e+03
2.16938384e+03	2.16925451e+03	2.16857990e+03	1.18826921e+05
2.16999131e+03	2.16787513e+03	2.16899130e+03	2.17068234e+03
1.18825497e+05	1.18826964e+05	2.16806527e+03	2.16867295e+03
2.17008695e+03	2.16803738e+03	2.16708830e+03	1.18825162e+05
2.16658882e+03	2.16627979e+03	2.16846368e+03	2.16781526e+03
2.16702554e+03	2.16946289e+03	2.16954914e+03	2.17109144e+03
2.16961113e+03	2.17009186e+03	2.16990630e+03	1.18827589e+05
1.18827266e+05	2.16822930e+03	2.16935024e+03	2.16912302e+03
1.18825851e+05	2.16956416e+03	2.17022601e+03	2.17009413e+03
2.17114685e+03	2.16918728e+03	2.16906738e+03	2.16879851e+03
1.18825129e+05	1.18826684e+05	2.16824461e+03	2.16787666e+03
2.16986038e+03	1.18827263e+05	2.16837901e+03	2.17044041e+03
2.16963167e+03	2.16906490e+03	2.17075757e+03	2.17110481e+03
2.17119573e+03	2.16979956e+03	2.16948012e+03	2.17035024e+03
2.16960022e+03	1.18826870e+05	1.18825401e+05	1.18826866e+05
2.17127274e+03	2.16996059e+03	2.16998505e+03	2.17028118e+03
1.18825159e+05	2.16895375e+03	2.16967438e+03	2.17092925e+03
2.16980103e+03	1.18826490e+05	1.18826944e+05	2.16882894e+03
2.17095943e+03	2.16962722e+03	2.16802887e+03	2.16841266e+03
2.16798328e+03	2.16789270e+03	2.16699305e+03	2.16863527e+03
2.17017040e+03	2.16974843e+03	2.16876126e+03	2.16804349e+03
2.16822345e+03	1.18824208e+05	2.16658182e+03	1.18824976e+05
2.16703640e+03	2.16816554e+03	2.16696932e+03	2.16902185e+03
2.17029609e+03	2.16752145e+03	2.17080398e+03	2.17020757e+03]

```
In [198]: 1 # Equation of the model
2
3 def predictLE4(weights, x):
4
5     equation = 0
6     for i in range(len(m4)):
7         equation += (weights[i]*x[i])
8     equation += intercept4
9     print("The Life Expectancy is " + str(equation))
```

Training a SGD Regression model:


```
In [199]: 1 # Fitting the Data
2 sgdr = SGDRegressor()
3 sgdr.fit(X_train, y_train)
4
5 # Predicting the Target Variable
6 y_pred5 = sgdr.predict(X_test)
```

Regression Plot

```
In [200]: 1 # Plotting the true and predicted value
2 fig = plt.figure(figsize =(20, 10))
3 sns.regplot(y_test, y_pred5)
4 plt.xlabel("LE Test")
5 plt.ylabel("LE Predicted")
6 plt.title("True Value vs Predicted Value")
7 plt.show()
```

C:\Users\joanj\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



```
In [201]: 1 # Chcking the accuracy of the model
2 print("The accuracy of the model is", sgdr.score(X_train, y_train))
```

The accuracy of the model is -8.436743678238087e+46

```
In [202]: 1 cv_score = cross_val_score(sgdr, X, y, cv = 10)
2 print("CV mean score: ", cv_score.mean())
```

CV mean score: -1.0706117933028202e+48

Boosting Efficiency

```
In [203]: 1 #generic function to fit model and return metrics for every algorithm
2 def boost_models(x):
3     #transforming target variable through quantile transformer
4     regr_trans = TransformedTargetRegressor(regressor=x, transformer=QuantileTransformer(output_distribution='normal'))
5     regr_trans.fit(X_train, y_train)
6     yhat = regr_trans.predict(X_test)
7     algoname= x.__class__.__name__
8     return algoname, round(r2_score(y_test, yhat),3), round(mean_absolute_error(y_test, yhat),2), round(np.sqrt(mean_squared_error(y_test, yhat)),2)
9
10 algo=[GradientBoostingRegressor(), lightgbm.LGBMRegressor(), xgboost.XGBRFRegressor()]
11 score=[]
12
13 for a in algo:
14     score.append(boost_models(a))
15
16 #Collate all scores in a table
17 pd.DataFrame(score, columns=['Model', 'Score', 'MAE', 'RMSE'])
```

Out[203]:

	Model	Score	MAE	RMSE
0	GradientBoostingRegressor	0.930	1.83	2.50
1	LGBMRegressor	0.955	1.37	2.00
2	XGBRFRegressor	0.924	1.86	2.61

FLASK Attempt:

```
In [74]: 1 import pickle
2 from joblib import dump, load
3 dump(linear_regressor, 'model.joblib')
```

Out[74]: ['model.joblib']

```
In [75]: 1 loaded_model = load('model.joblib')
```

```
In [76]: 1 # !pip3 install kaleido
```

```
In [77]: 1 def make_picture2(training_data, model, new_input_arr, output_file='predictions_pic.svg'):
2
3     # x_new = [1,2,3,4,5,6,7,8,9,10]
4     x_new=np.arange(10).reshape((10, 1))
5     preds = model.predict(np.array(x_new).reshape(1,-1))
6     fig = px.scatter(x=y_test, y=y_pred, title="Predicted Life Expectancy", labels={'x': 'True Values',
7                                         'y': 'Predicted Value Life Expectancy'})
8     x_new=np.array(x_new)
9
10    fig.add_trace(
11        go.Scatter(x=x_new.reshape(x_new.shape[0]), y=preds, mode='lines', name='Model'))
12
13    if new_input_arr is not False:
14        new_preds = model.predict(np.array(new_input_arr).reshape(1,-1))
15        new_input_arr = np.array(new_input_arr)
16        fig.add_trace(
17            go.Scatter(x=new_input_arr.reshape(new_input_arr.shape[0]), y=new_preds, name='New Outputs', mode='markers', marker=dict(
18                color='purple',
19                size=10,
20                line=dict(
21                    color='purple',
22                    width=2
23                )))
24    # fig.write_image(output_file, width=800)
25    return fig
```

```
In [78]: 1 def floats_string_to_np_arr(floats_str):
2     def is_float(s):
3         try:
4             float(s)
5             return True
6         except:
7             return False
8     floats = [float(x) for x in floats_str.split(",") if is_float(x)]
9     return floats
10    # floats.reshape(len(np.array(floats)),1)
```

```
In [79]: 1 import pickle
```

```
In [80]: 1 pickle.dump(linear_regressor, open('pickle_file', 'wb'))
```

